



Secret AI Rails – Pilot Case Studies

Overview

This document profiles our pilot customers and design partners who are validating Secret AI Rails in production environments. These early adopters demonstrate market demand and provide critical feedback for product development.

Pilot Program Summary

Metric	Status
Active Pilots	2-3 SMB customers
Design Partners	MIT (academic/research)
Pilot Duration	3-6 months
Focus Areas	Business communications, document processing, customer operations

Case Study Template

Use this template for documenting each pilot engagement

[Company Name]

Industry:

Company Size:

Location:

Pilot Start Date:

Champion/Contact:

The Challenge

[What problem were they trying to solve?]

Why They Chose Secret AI Rails

[What attracted them to our platform?]

Solution Deployed

[Which workflows are they running?]

Implementation Timeline

[How long to get to production?]

Results (Expected/Actual)

Metric	Before	After	Change

Customer Quote

> "[Testimonial from customer champion]"

> – Name, Title

Key Learnings

[What did we learn from this pilot?]

Expansion Potential

[Where could they grow usage?]

Pilot #1: [SMB – Professional Services]

Company Profile

Attribute	Details
Industry	Professional Services / Consulting
Company Size	25-50 employees
Location	United States
Pilot Start Date	[Month Year]
Champion	Head of Operations

The Challenge

This professional services firm was struggling with:

- **Email overload:** 200+ client emails daily requiring triage and response
- **Inconsistent response times:** Ranging from 2 hours to 2 days depending on staff availability
- **No audit trail:** Difficulty tracking what was communicated to which client
- **Compliance concerns:** Clients in regulated industries requiring proof of communications

Manual process cost: ~\$8,000/month in staff time for email management alone

Why They Chose Secret AI Rails

1. **Confidentiality:** Client data stays encrypted, even during AI processing
2. **Audit trail:** Every AI action is logged and provable
3. **Human gates:** Sensitive communications require approval before sending
4. **Existing relationship:** Already using our confidential VM infrastructure

Solution Deployed

flowchart LR subgraph Trigger Email[Incoming Email] end subgraph AIRails["Secret AI Rails"]

Classify[Classify Intent] Extract[Extract Key Info] Draft[Draft Response]
 Gate{Approval
Required?} Human[Human Review] Send[Send Response] Log[Log to
 Proof Ledger] end subgraph Output Reply[Email Sent] Audit[Audit Trail] end Email --> Classify
 --> Extract --> Draft --> Gate Gate -->|High Risk| Human --> Send Gate -->|Low Risk| Send
 Send --> Reply Send --> Log --> Audit

Workflows Deployed:

- Email classification and routing
- Automated response drafting
- Client inquiry triage
- Meeting scheduling coordination

Implementation Timeline

Phase	Duration	Activities
Discovery	1 week	Requirements gathering, workflow mapping
Configuration	2 weeks	Workflow setup, policy configuration, integration
Testing	1 week	Dry runs with synthetic data, edge case testing
Pilot Launch	Ongoing	Production deployment with monitoring

Total time to production: 4 weeks

Results

Metric	Before	After (Projected)	Change
Avg. Response Time	8 hours	15 minutes	-94%
Emails Handled/Day	200	200	—
Staff Hours on Email	6 hrs/day	1 hr/day	-83%
Monthly Cost	\$8,000	\$2,500	-69%
Response Consistency	Variable	95%+	Standardized

Metric	Before	After (Projected)	Change
Audit Compliance	Manual	Automatic	100% coverage

Customer Quote

"We were drowning in client emails. Now AI handles the routine stuff instantly, and I only see the messages that actually need my attention. The audit trail alone is worth it — our regulated clients love that we can prove exactly what was communicated."

— [Name], Head of Operations

Key Learnings

- Policy tuning is critical:** Initial policies were too conservative; iterative refinement improved automation rate from 40% to 75%
- Proof ledger is a selling point:** Customer uses audit exports in client reporting
- Integration simplicity matters:** Gmail OAuth setup needed to be streamlined

Expansion Potential

- Document intake and processing (contracts, proposals)
- Slack integration for internal coordination
- Client portal with self-service status checks
- Voice call transcription and follow-up

Pilot #2: [SMB – E-Commerce / Retail]

Company Profile

Attribute	Details
Industry	E-Commerce / Direct-to-Consumer Retail
Company Size	15-30 employees

Attribute	Details
Location	United States
Pilot Start Date	[Month Year]
Champion	Customer Experience Manager

The Challenge

This growing e-commerce brand faced:

- **Customer support volume:** 150+ support tickets/day, scaling with growth
- **Repetitive inquiries:** 70% of tickets were order status, returns, basic product questions
- **Limited hours:** Support only available during business hours
- **No knowledge base integration:** Agents manually searching for product information

Pain point: Hiring more support staff wasn't sustainable; needed automation that customers could trust

Why They Chose Secret AI Rails

1. **24/7 availability:** AI handles inquiries around the clock
2. **Knowledge-grounded responses:** Answers based on actual product data, not hallucinations
3. **Escalation gates:** Complex issues route to humans automatically
4. **Proof for disputes:** Complete record of every customer interaction

Solution Deployed

```

flowchart TB
    subgraph Inputs
        Ticket[Support Ticket]
        Chat[Live Chat]
        Email[Customer Email]
    end
    subgraph AIRails ["Secret AI Rails"]
        Classify[Classify Issue Type]
        KB[Query Knowledge Base]
    end
    subgraph Routing
        Simple{Simple<br/>Query?}
        Auto[Auto-Respond]
        Complex[Route to Human]
    end
    Respond[Generate Response]
    Proof[Log to Ledger]
    subgraph Outputs
        Resolution[Customer Resolution]
        Escalation[Human Agent Queue]
        Audit[Audit Trail]
    end
    Inputs --> Classify
    Classify --> KB
    KB --> Simple
    Simple -->|Yes| Auto
    Auto --> Respond
    Respond -->|No| Complex
    Complex --> Escalation
    Escalation --> Resolution
    Escalation --> Audit
    Resolution --> Proof
    Proof --> Audit

```

Workflows Deployed:

- Product Q&A chatbot (knowledge-grounded)
- Order status automation
- Return/refund request processing (with approval gates)
- FAQ response automation

Implementation Timeline

Phase	Duration	Activities
Discovery	1 week	Ticket analysis, FAQ identification
Knowledge Ingestion	1 week	Product catalog, policies, common answers
Configuration	2 weeks	Workflow setup, response templates, policies
Testing	1 week	A/B testing with subset of tickets
Pilot Launch	Ongoing	Full deployment with human oversight

Total time to production: 5 weeks

Results

Metric	Before	After (Projected)	Change
Tickets Handled/Day	150	150	—
Auto-Resolution Rate	0%	65%	+65%
Avg. First Response	4 hours	2 minutes	-99%
Support Hours/Day	8 hrs	3 hrs	-63%
CSAT Score	3.8/5	4.2/5	+10%
Cost per Ticket	\$4.50	\$1.80	-60%

Customer Quote

"Our customers don't care if it's AI or human — they care about fast, accurate answers. Secret AI Rails gives us both, and I can prove to my boss exactly what's happening with

every interaction."

— [Name], Customer Experience Manager

Key Learnings

1. **Knowledge base quality is everything:** Garbage in, garbage out — invested heavily in KB curation
2. **Escalation thresholds need tuning:** Started conservative, gradually increased automation
3. **Customers appreciate transparency:** "AI-assisted response" disclosure didn't hurt CSAT

Expansion Potential

- Social media monitoring and response
- Proactive order issue detection
- Review response automation
- Voice support (STT-TTS)

Pilot #3: [MIT – Academic/Research Partnership]

Institution Profile

Attribute	Details
Institution	Massachusetts Institute of Technology
Department	[Relevant Department/Lab]
Location	Cambridge, MA
Partnership Start	[Month Year]
Champion	[Professor/Researcher Name]

The Opportunity

MIT represents a strategic design partner for several reasons:

1. **Research credibility:** Academic validation of confidential AI execution
2. **Cutting-edge use cases:** Research workflows with sensitive data requirements
3. **Talent pipeline:** Access to top AI/security talent
4. **Publication potential:** Joint research on trusted AI execution

Partnership Focus Areas

Area	Description
Research Data Processing	AI workflows on sensitive research data with confidentiality guarantees
Administrative Automation	Grant management, scheduling, correspondence
Verification Research	Academic study of proof-of-work systems for AI actions
Security Analysis	Formal verification of trust model

Potential Research Collaboration

flowchart TB
MIT["MIT Research"] --> Research[Research Questions]
Data[Sensitive Datasets] --> Validation[Validation & Testing]
Publication[Academic Publication] --> end
subgraph SecretAI["Secret AI"]
Platform[AI Rails Platform] CC[Confidential Compute] Proof[Proof Systems]
end
subgraph Outcomes
Papers[Published Research] Product[Product Improvements]
Credibility[Market Credibility]
end
Research --> Platform
Data --> CC
Validation --> Proof
Platform --> Papers
CC --> Product
Proof --> Credibility
Publication --> Papers

Value Exchange

MIT Receives	Secret AI Receives
Confidential AI infrastructure for research	Academic validation and credibility

MIT Receives	Secret AI Receives
Early access to platform features	Research insights and feedback
Potential research funding collaboration	Publication co-authorship
Student project opportunities	Talent pipeline

Expected Outcomes

1. **Joint publication** on confidential AI execution frameworks
2. **Case study** on research data processing with trust guarantees
3. **Security audit** findings and recommendations
4. **Reference customer** for enterprise sales conversations

Quote

"The ability to run AI workflows on sensitive research data while maintaining confidentiality and producing verifiable audit trails opens new possibilities for responsible AI in research."

— [Name], [Title], MIT

Pilot Success Metrics (Aggregate)

Across All Pilots

Metric	Target	Current Status
Pilots Active	3	2-3
Time to First Value	<6 weeks	4-5 weeks avg
Workflow Automation Rate	>50%	60-75%
Customer Satisfaction	>4.0/5	4.1/5 avg

Metric	Target	Current Status
Zero Security Incidents	0	0
Proof Ledger Accuracy	100%	100%

Pilot-to-Paid Conversion

Stage	Count	Conversion
Pilot Started	3	—
Pilot Completed	—	TBD
Converted to Paid	—	Target: 67%+
Expanded Usage	—	Target: 50%+

Learnings & Product Implications

What's Working

Finding	Product Implication
Proof ledger is a major selling point	Emphasize audit capabilities in marketing
Policy tuning requires iteration	Build better policy templates and recommendations
Integration setup is friction point	Invest in OAuth flows and guided setup
Knowledge base quality drives success	Add KB quality scoring and recommendations

What Needs Improvement

Finding	Action
Initial policy configuration is complex	Create industry-specific templates
Dashboard needs more real-time visibility	Prioritize live metrics in v1
Escalation routing needs refinement	Add ML-based routing recommendations

Feature Requests from Pilots

Request	Priority	Status
Slack integration	High	In v1 scope
Document attachment processing	High	In v1 scope
Custom reporting/exports	Medium	Planned
Multi-language support	Medium	Backlog
Voice integration	Medium	In v1 scope (basic)

Testimonials Summary

"The audit trail alone is worth it — our regulated clients love that we can prove exactly what was communicated."

"Our customers don't care if it's AI or human — they care about fast, accurate answers."

"We were drowning in client emails. Now AI handles the routine stuff instantly."

"The ability to run AI workflows on sensitive research data while maintaining confidentiality opens new possibilities."

Next Steps

Pilot Expansion Plan

Timeline	Action
Current	Complete v1 workflows with existing pilots
Month 2-3	Gather quantitative results and testimonials
Month 3-4	Publish case studies with permission
Month 4-6	Expand to 5-10 additional pilots
Month 6+	Convert pilots to paid customers

Target Pilot Profiles (Next Wave)

Segment	Use Case	Why
Legal Services	Document processing, client communications	High compliance need
Healthcare Admin	Patient communications, scheduling	HIPAA requirements
Financial Services	Client reporting, compliance workflows	Audit requirements
SaaS Companies	Customer support, onboarding	Scale challenges

Appendix: Pilot Agreement Template

Key Terms

Term	Standard
Duration	3-6 months
Pricing	Discounted or free during pilot
Commitment	Dedicated champion, weekly feedback
Success Criteria	Defined upfront, measured monthly
Case Study Rights	Permission to publish (anonymized if needed)
Conversion Terms	Preferred pricing for pilot-to-paid

Pilot Checklist

- Pilot agreement signed
 - Success criteria defined
 - Champion identified
 - Technical requirements confirmed
 - Integration access granted
 - Kickoff meeting completed
 - Weekly check-in scheduled
 - Feedback channel established
-

Document Version: 1.0

Last Updated: Jan 20, 2026

Status: Active Pilots in Progress